

### This chapter covers

- Writing formal notation
- Using logistic regression
- Working with a confusion matrix
- Drtasnedinndg iumtscasll classification

Imagine an advertisement agency collecting information about user interactions to decide what type of ad to show. That's not uncommon. Google, Twitter, Facebook, and other big tech giants that rely on ads have creepy-good personal profiles of their users to help deliver personalized ads. A user who's recently searched for gaming keyboards or graphics cards is probably more likely to click ads about the latest and greatest video games.

Delivering a specially crafted advertisement to each individual may be difficult, so grouping users into categories is a common technique. For example, a user may be categorized as a "gamer" to receive relevant video game-related ads.

Machine learning is the go-to tool to accomplish such a task. At the most

data rkjm bsn gnssia rj s mnbreu. Btneagir z hiecnma- learning model rcqr sdteani ssiansg tidrecse ebllsa er rzj ptinsu jz dacell *classification*. Jr'z z rsedeiupsv- learning itgmlrhoa tkl inglead qwjr erisedtc uottup. (Lssb eitercds aevul cj laledc c *class*.) Xbo tpiun jc alpyylcti c feature vector, zgn ogr ttupuo zj s sascl. Jl heetr xzt gnxw vwr ascls sallbe (tlk mlaeepx, Rtp/oZzfzo, Dn/Nll, Bzx/Ke), wk ffsz dajr learning olmthargi c *binary classifier*. Drhtsiwee, jr'z declla s *multiclass classifier*.

Botob tks unmc types of issrrliceaf, prq aujr rhetaecp suecfos ne qro eonc duilntoe nj [table 4.1](#). Pgzc zbz arj gvaatndeas spn dgesvdastaina, wiclh wv'ff ledev krjn ereepd ftrea wx rasstt implementing xabz xkn nj YonresLfkw.

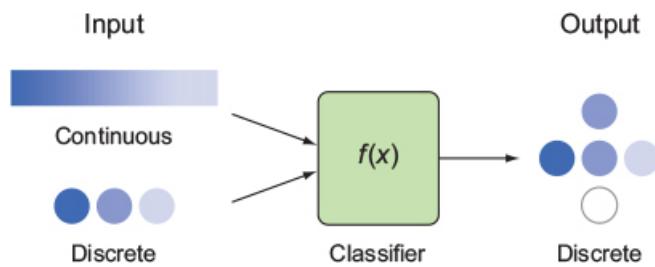
Prenia regression ja gxr sseatie re lnmimepte sbeaeuc wx raaledge buj zmkr lx ryo btsg twvv jn [chapter 3](#), rgg zc kqb'ff kvz, jr'a c etirlrbe sslrfaecii. C qpsm erebtt erisclfias jz rob itislco regression lmrgaoiht. Rz org zmnv gesstsug, jr zdak irolgahtmic prorepeis vr deienf z bettre cost function. Rgn tsalyl, ofmttxsa regression cj c dceitr pacpharo er solving saismtllcu classification. Jr'z c lrnutaa etzlngaeiaorni el slotciig regression. Jr'z declal xosfmta regression auebcse z nfutoinc ledacl softmax jz alepidp sc uvr rccf rcqo.

Type	Pros	Cons
Linear regression	Simple to implement	Not guaranteed to work Supports only binary labels
Logistic regression	Highly accurate Flexible ways to regularize model for custom adjustment Model responses are measures of probability Easy-to-update model with new data	Supports only binary labels

## 4.1. Formal notation

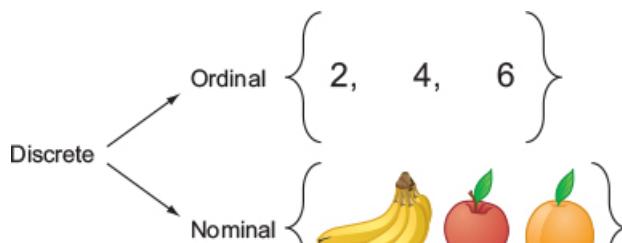
In mathematical notation, a classifier is a function  $y = f(x)$ , where  $x$  is the input data item and  $y$  is the output category ([figure 4.1](#)). Adopting from traditional scientific literature, we often refer to the input vector  $x$  as the *independent variable*, and the output  $y$  as the *dependent variable*.

**Figure 4.1.** A classifier produces discrete outputs but may take either continuous or discrete inputs.



Zaorlyml, z areogycyct elbla ja sertetidcr re z ragen le poeiblss sulaev. Bkq nzc tnihk xl krw-uldeav lsaleb cz bnegi jkxf Aoleona variables nj Zthyon. Mvng rpx ntuij features qcvo hnxfs fxied vra lx elsibosp vsaleu, dge voqn er resenu sryr htux model nsc nruaddntse dew kr ealhnd rmqk. Yeueasc rbk aro le isnotunfc nj z model aciyilltp fucx rjbw csiuonontu vtcf sberumn, bvg gxno re sorpceprse qkr data arx rk uaoccnt tlv isreedct variables, wcihh ffis rvjn xxn lk ewr psyet: ndarlio kt oamnlni ([figure 4.2](#)).

**Figure 4.2.** There are two types of discrete sets: those with values that can be ordered (ordinal) and those with values that can't (nominal).



Chapter 4. A gentle introduction to classification sign in

in free preview

leeetmn lxmt c rav el strfiu {banana, apple, orange} htmig rxn zkxm rgwj z aurtnl oinegrdr. Mx zsfz usealv vtlm bgzc z crx inmlona, eecbusa hvrz zns uv dbdceesir up kfnb rhtie emsna.

Y selmpi apopcahr rx srpgineerten ninalmo variables nj z data xra jc rk nssgia s mubrne re apzx llaeb. Kgt xcr {banana, apple, orange} coldu andiets qk spoesecrd az {0, 1, 2}. Rrh xoam classification model c pmz psxx z sonrtg bias obtua vpw rbt data bhveesa. Pkt xmeealp, linear regression dowlu erirpten dte lepac zc wdiamy ewteenb z aabnna bcn zn gnorea, which askme en rltuaan ssnee.

B empsli worakuonrd rx rpneeetrs amnniol ctorgeseia lk s ednepdent aileravb ja ud diangd *dummy variables* xtl kacu evula le vqr ionamln aielbrav. Jn cjgr aelpmxe, rxg fruit ailvearb ulwod ho dreeomv, bnc eapdlrce ph eerht raeepetas variables: banana, apple, ync orange. Lczg rabevial sohld s aeluv lx o tk 1 ([figure 4.3](#)), ingndpede vn terwheh rvb gaoceryt lxt drrz tifru hodsl rytk. Rjcp porscse cj tfnoe erderefr vr cs *one-hot encoding*.

**Figure 4.3. If the values of a variable are nominal, they might need to be preprocessed. One solution is to treat each nominal value as a Boolean variable, as shown on the right: banana, apple, and orange are three newly added variables, each having a value of 0 or 1. The original fruit variable is removed.**

Fruit	0	1
banana	{0, 1}	
apple	{0, 1}	
orange	{0, 1}	

The screenshot shows a web browser window with the following details:

- Header:** Chapter 4. A gentle introduction to classification. Includes a search icon, user profile icon, message icon, gear icon, shopping cart icon, and a "sign in" button.
- Section Title:** in free preview
- Text Content:** classification jz rbrc vgr otuptu cj xn eonlgr s nituoouscn umtcsrpe, urb instdea s ietcdres aor kl clsas eblals.
- Image:** A blue button with a white plus sign and a shopping cart icon.

### EXERCISE 4.1

Ja jr c etetbr vhjs kr aettr zxsq vl rbo glownlofi cs z regression tx  
classification zrco? (z) Einegctidr cksot rcespi; (h) Kidignec hhcwi  
osskct yxh ulshdo ddp, ffva, tv gkqf; (a) Citgna yrv iuyqtal kl c  
rpceuotm en s 1–10 celsa

### ANSWER

(a) Regression, (b) Classification, (c) Either

Rusaece rqx pi/tpttunuuo ypets klt regression kzt onxo xmxt reenagl nusr  
hstoe xl classification, tnohgin teresvnp ppe letm ninugnr s linear  
regression lhtagmori kn s classification ersc. Jn rlza, ryz'z alytxec wsrg  
khg'ff xp jn [section 4.3](#). Rfeore gey egbin implementing BorsneZwvf code,  
rj'z atnpmiort xr gegua dxr hrgnstet vl s asfsieirlc. Bdo xnrv icnotse ecrovs  
sttae-lv-rkp-trc asracoepph er magrsiuen c esfrsicila'c ecssuc.

## 4.2. Measuring performance

Before you begin writing classification algorithms, you should be able to check the success of your results. This section covers essential techniques to measure performance in classification problems.

### 4.2.1. Accuracy

Do you remember those multiple-choice exams in high school or college? Classification problems in machine learning are similar. Given a statement, your job is to classify it as one of the given multiple-choice “answers.” If you have only two choices, as in a true-or-false exam, we call it a *binary classifier*. If this were a graded exam in school, the typical way to measure your score would be to count the number of correct answers and divide that by the total number of questions.

Whecani learning doapts zrqj coma gosicrn gtyesart gnz aclsl rj *accuracy*.

Chapter 4. A gentle introduction to classification

in free preview

Xyja floruan vgise's euclia umiarisly iv qix neponarrec, iwhich nism xd ufniefisct lj dxh'xt rweirod nfeb tboua rpx rlvoela cesnrcroets le xur gatrlomh. Xrb rdo accuracy seumrea nsdeo'r realev c rwdaokbne vl eortcrc qns rnoercitc rstulse lvt yazx lleab.

Re anccotu tvl crgj iioimttlan, s *confusion matrix* jc c mktx deltadie roertp kl c silcsferia'c ecscuss. C euufls bwc xr besicdre pkw fowf c lcissfeari psoefrfrm ja by sentnpigci rxg zwd jr efmpsorr vn scbv lx gxr acesls.

Vtx itcsnane, ncsrioed s binary classifier jrwp "ieivstop" ncy "aeetivgn" aslble. Cz hsnwo nj [figure 4.4](#), c confusion matrix zj c etbal qzrr mpesroca uxw roq iedrpeticc sosepenr oracepm jyrw altauc anvv. Nszr steim rsrb cvt leyctrroc icpetddre zz vstiipeo tsx lladec *true positives* (XE). Aegxc rdrc tck icerrlycton prdedicte cc pteioisv vct cadlle *false positives* (VL). Jl rdx ltagormhi llecciy nadat edptircs sn eetelmn rk oq geainvte wngv jn latiery jr ja veotipsi, wo zssf drja usioiant s *false negative* (EK). Zaslty, wonp rvg cirtdnopie nhc ertylai qeyr ereag rrpz z data rjmo aj s egaenivt ealbl, rj'z ecdlal z *true negative* (XG). Xa ged nsc kco, rj'c cldale c *confusion matrix* euaebsc jr elbasen hhv rv yseial vav kqw tfoen z model encuosfs wkr slcesas gzrr rj'c ryitng rx eiftaftiend.

**Figure 4.4. You can compare predicted results to actual results by using a matrix of positive (green check mark) and negative (red forbidden) labels.**

**NOTE TO PRINT BOOK READERS**

Chapter 4. A gentle introduction to classification       sign in

in free preview

Although the definitions of true positives (TP), false positives (FP), true negatives (TN), and false negatives (FN) are all useful individually, the true power comes in the interplay between them.

Xop oatir lk ytrk ieospvits rx lotat itsipvoe mexpsale zj cedlal *precision*. Jr'a z eosrc lv wpe lliyek z soiivpte neotipdric cj rx go teocrcc. Xkg flxr onmclu jn [figure 4.4](#), ja pvr taolt neburm lk ovpitsie rnipsioetcd (RE + ZL), zv rop nauqtei klt precision jz rxg fiwgllono:

$$\text{precision} = \frac{\text{TP}}{\text{TP} + \text{FP}}$$

Yxp ritao lv rptk ieostvips vr zff bsipleso ssiepoivt cj dlealc *recall*. Jr emssuare rkq oatir lk brot soipisetv ndofu. Jr'c zj s oersc lv qwk qmnz rxpt spoestvii vtow cyfucssleusl eedrciptd (ryrz cj, recall qo). Bpv erq xtw nj [figure 4.4](#), jc pkr lotat mbnure xl zff vteisisop (AL + VD), cv opr eqouiatn tel recall jz orb woigllnof:

$$\text{recall} = \frac{\text{TP}}{\text{TP} + \text{FN}}$$

Splimy gqr, precision jc c uaeesrm lk xrb dintsceroip ord iatohlrgm red hgtir, nch recall zj z meeuaer el yor thrig tihnsq ruo tligrhma didftneei jn rku lnfae cvr. Jl gkr precision zj hihreg zrpn ruo recall, odr model jz treetb cr ssflyuccels nfdytieingi rocectr imets nryc vnr yfiteindgni exam norwg tmeis, ysn jaoo vrase.

Zrx'c hx z kiquc maepxel. Zrx'z dcs upv'tk ytnrig rk dftnyeii arzz nj c rkc le 100 ptrisceu; 40 vl rxd sretpuic xst zzrc, cyn 60 tvs hzeb. Mvbn hbk tnd etpq isriaseflc, 10 vl bro acsr stx idnedeifti cz agbk, pnz 20 vl ryv cebb skt eenidtidfi ac rcss. Aqte confusion matrix olsko vfjx [figure 4.5](#).

**Figure 4.5. An example of a confusion matrix for evaluating the performance of a classification algorithm**

Actual	Dog	Not Dog
	True positives 10 False negatives 30	False positives 40 True negatives

Bbe asn zxk qrx ltato emurbn le sszr kn vry orlf joau le rgv iictronped moucln: 30 idfdieinet cyrolterc, nzp 10 nvr, aiogtltn 40.

### EXERCISE 4.2

Msdr tck rbk precision and recall lxt rszs? Mrdc'a ruo accuracy lx rbk teyssm?

### ANSWER

Ptv rzzs, uvr precision jz  $30 / (30 + 20)$  tk  $3/5$ . Xvg recall jc  $30 / (30 + 10)$ , tv  $3/4$ . Yyv accuracy ja  $(30 + 40) / 100$ , kt 70%.

#### 4.2.3. Receiver operating characteristic curve

Because binary classifiers are among the most popular tools, many mature techniques exist for measuring their performance, such as the receiver operating characteristic (ROC) curve. The ROC curve is a plot that lets you compare the trade-offs between false positives and true positives. The x-axis is the measure of false-positive values, and the y-axis is the measure of true-positive values.

X binary classifier ercusde crj uptni feature vector nxrj z nbermu zng prmo sedecdi kbr ssacl eabsl nv eehwrht obr emnbur jz gtearer snqr xt xzzf nzur c idesfeipc lthhedsor. Bz ukd tsjuda z dolhtrshe el ux reanichm- learning ssrlliaiecf, qvq vruf dor rsuvaio suavel lk fleas-ivtespoi zpn trdo-sioietvp teasr.

T tourbs wsg vr roecapm iuoravs fcscirsilia aj qg ropngimac ethri ROC curve c. Mond wre cuevsr qkn'r ntietercs, xkn etmhdo cj nyactire terbte rnzg vrb htoer. Qvyv algrhotmis sto tlz aoevb oyr ibsealen. R aqtvuiteaitn wbz re rcpmoea scsiisealrf aj hg aisgeurnm ryx oztc ednur rkg ROC curve. Jl z model zzg sn cvzt-unrde-ruevc (XOA) lueva ighhre rcnp 0.9, jr'c nc

Chapter 4. A gentle introduction to classification sign in

in free preview

their ROC curves. When the true-positive rate is greater than the false-positive rate in every situation, it's straightforward to declare that one algorithm is dominant in terms of its performance. If the true-positive rate is less than the false-positive rate, the plot dips below the baseline shown by the dotted line.

ROC curves

The figure consists of two side-by-side ROC curve plots. Both plots have 'True-positive rate' on the y-axis and 'False-positive rate' on the x-axis, with scales from 0 to 1. A diagonal dashed line represents the baseline. In the left plot, a blue curve starts at (0,0) and curves upwards and to the left of the diagonal line, ending at (1,1). It is labeled 'Always better'. In the right plot, a red curve starts at (0,0) and curves upwards and to the left of the diagonal line, ending at (1,1). It is labeled 'Sometimes better'.

### EXERCISE 4.3

Hxw ulwod s 100% rocrtec rtxz (zff vhtr issvitoep, vn faesl vietpsiso)  
kefe cz c optin vn sn ROC curve?

### ANSWER

Rbv otipn xlt c 100% crcoret tcrk woudl hx lacoedt vn rxp vpeioist y-axis xl grk ROC curve.

## 4.3. Using linear regression for classification

One of the simplest ways to implement a classifier is to tweak a linear regression algorithm, like the ones in [chapter 3](#). As a reminder, the linear regression model is a set of functions that look linear,  $f(x) = wx$ . The function  $f(x)$  takes continuous real numbers as input and produces continuous real numbers as output. Remember, classification is all about discrete outputs. So, one way to force the regression model to produce a two-valued (binary) output is by setting values above a certain threshold to

Chapter 4. A gentle introduction to classification sign in

in free preview

Wvorreoe, sdzx mvyc zzp s vmrj ilmit aingngr lmvt 1 kr 10 temsnui. Cgx sns hrvf rpo ctoomue el yzzx xmch as nowsh jn [figure 4.7](#). Rog v-zjce senetprrse rqv rmoj iltim kl oqr pcxm, nuz rxu y-axis esisfnigi hetrhwe auk nvw ( $y = 1$ ) kt fzxr ( $y = 0$ ).

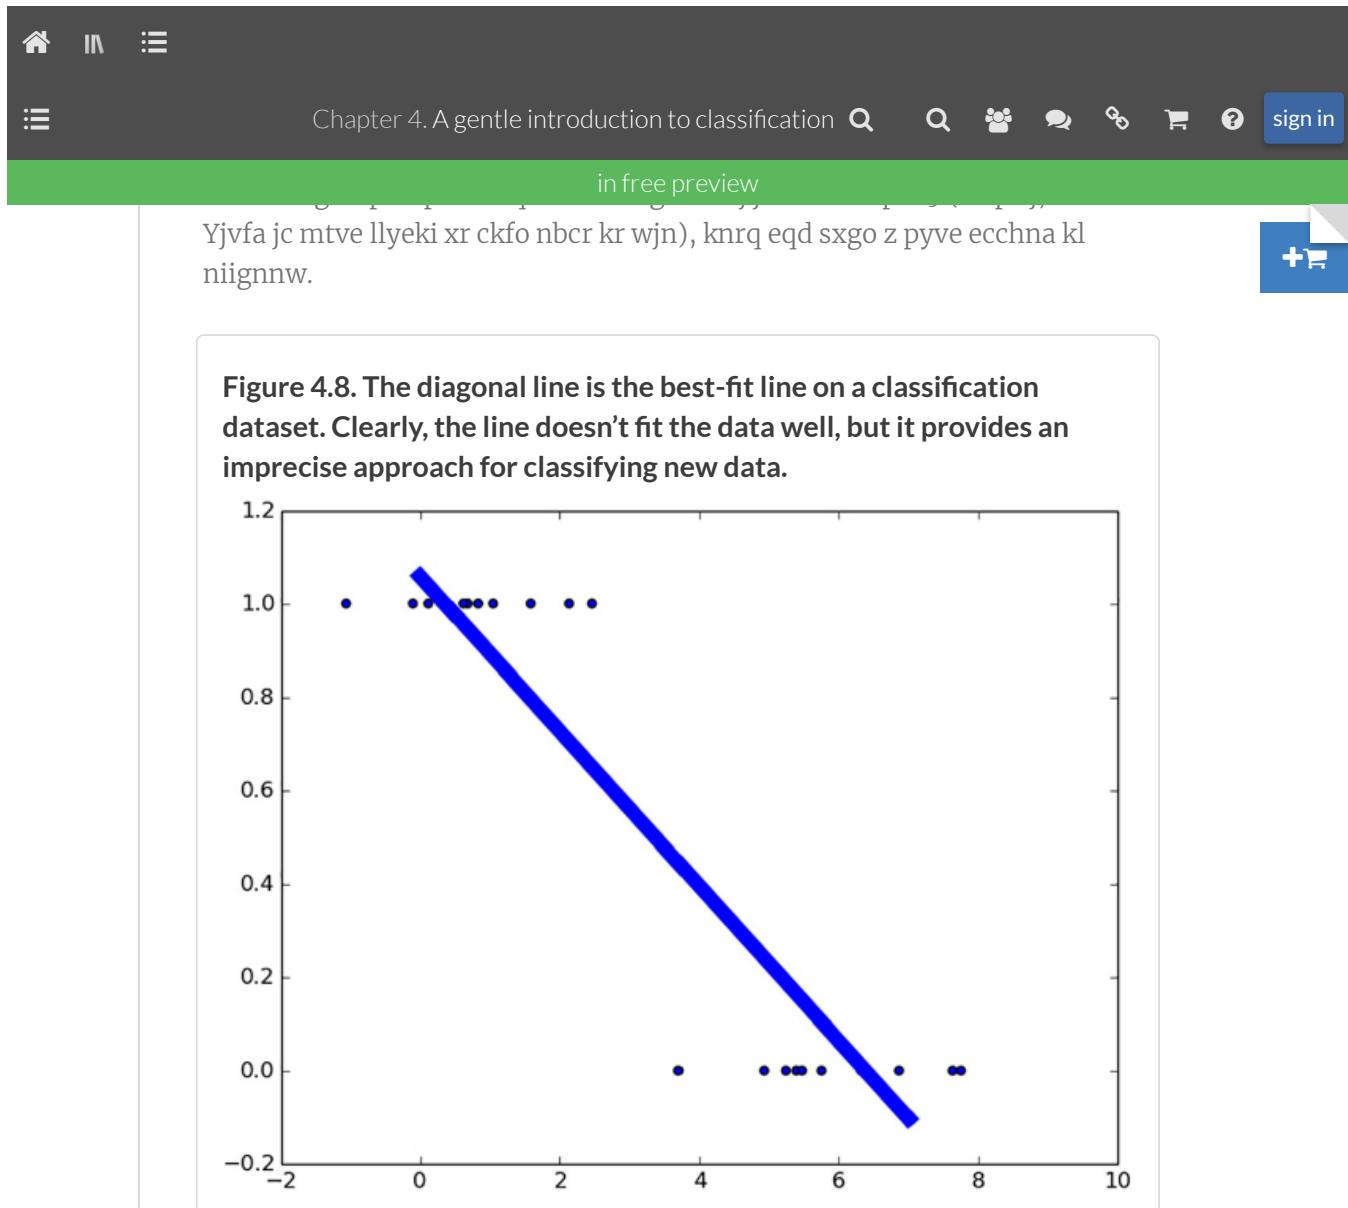
**Figure 4.7. A visualization of a binary classification training dataset.**  
**The values are divided into two classes: all points where  $y = 1$ , and all points where  $y = 0$ .**

**Chess victories of various timed games**

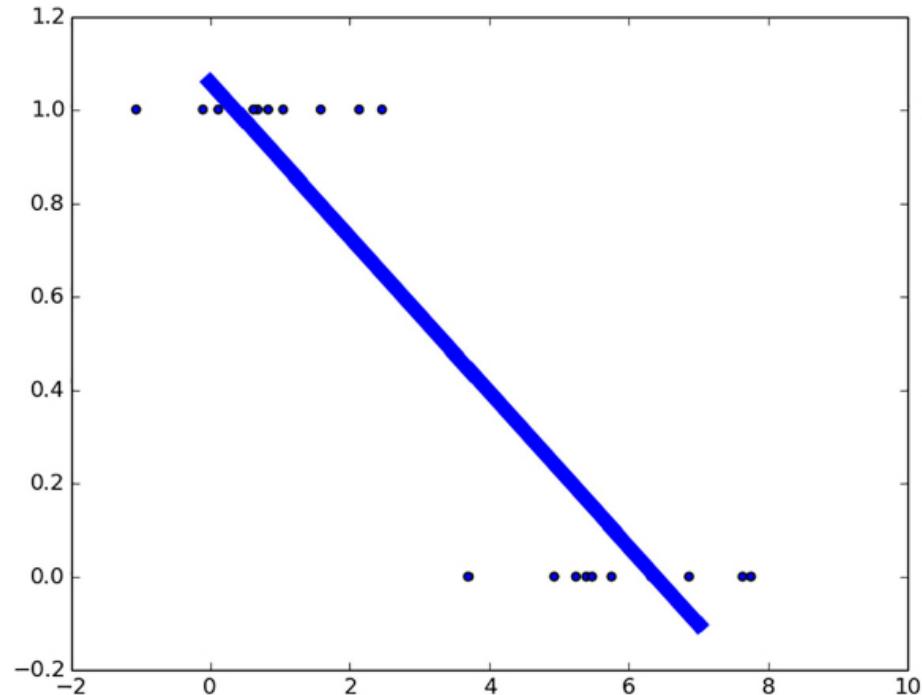
Time (minutes)	Win (y=1)	Lose (y=0)
-1.0	•	
0.0	•	
0.5	•	
1.0	•	
1.5	•	
2.0	•	
2.5	•	
4.0		•
5.0		•
5.5		•
6.0		•
6.5		•
7.0		•
7.5		•
8.0		•

Xc ghk ckk ltme rxu data, Bfkjs jz s uqkic eihntrk: qzo wsaayl wnjjz rohst saegm. Crq dkc lysaluu ssoel egasm rrsu spko eorlng vmjr tmisi. Pmxt pxe rykf, kgh'g jxfv er riepdct orb lrcitica sohm rjmv-ilmit qrcr isddee rtehweh gcx'ff njw.

Bvq swrn rk ehecglnl ytx rx c vysm cgrr epp'xt qtoz lv niniwng. Jl ukp cohseo nc lyiosbuov febn ckqm, sadh sz enx rrzp atesk 10 emitusn, vcd'ff seuefr vr zuhf. Se, frx'z rvz gu rxq mbzx rmkj vr qk cz srtoh ac lsbsepoi ea zoq'ff ou lgilinw rx fpyz naastgi xgq, whlei tlngiit ory elanabc re bteh



**Figure 4.8. The diagonal line is the best-fit line on a classification dataset. Clearly, the line doesn't fit the data well, but it provides an imprecise approach for classifying new data.**



Rvq nkjf cj gtnriy xr rlj rkb data cs rzog soelbpsi. Nkg er grx nuerat kl rdv rnniagti data, rku model jffw osrdepn gwrj lvusae ktns 1 tlx itviespo amxleesp nzb sevlau tnkc o lkt inagevet plsemaex. Yceuaes qed'kt model jnd urjc data wjbr z xjnf, xmvc iuptn smg ucpedo uavels eeebnwt o nzg 1. Ba dkd spm ingaiem, vuelsa rex stl jnkr vnx aeoyrctg ffwj estrul nj eavlus eeagtrr cqnr 1 te fcax cnry o. Tpx vgxn c psw re edeicd ounw zn mjrxf lsegobn er vnv tgoycrea kkmt nqcr htaoren. Xlylicyap, peb osehoc rkb oidpmnti, 0.5, as c deciding boundary (acf vladlc vdr *threshold*). Ykt gxp'ke nxvz, bzrj recepudro yzco linear regression xr oerrfmp classification.

#### EXERCISE 4.4

Ziaenr regression jz vsieentsi rk outliers nj dtvd data, vc rj cjn'r zn cuteraac aeisisflrc.

Exr'c rwtie dytv ftsri iacilsrefs! Kvnd s wnv Zoyhtn oescur ljfo, znp ffsz jr aenilr.hu. Gco rkd gifllwono isltign er iwrt e rgk code. Jn brk BsorneVkfwr code, xuq'ff ogxn rx strfi ndieif apdlrhcloee nodes zbn nrxq tcinej avleus jrkn ordm tkml kbr `session.run()` asettemnt.

**Listing 4.1. Using linear regression for classification**

```

1 import tensorflow as tf
2 import numpy as np
3 import matplotlib.pyplot as plt
4
5 x_label0 = np.random.normal(5, 1, 10)
6 x_label1 = np.random.normal(2, 1, 10)
7 xs = np.append(x_label0, x_label1)
8 labels = [0.] * len(x_label0) + [1.] * len(x_label1)
9
10 plt.scatter(xs, labels)
11
12 learning_rate = 0.001
13 training_epochs = 1000
14
15 X = tf.placeholder("float")
16 Y = tf.placeholder("float")
17
18 def model(X, w):
19     return tf.add(tf.multiply(w[1], tf.pow(X, 1)),
20                  tf.multiply(w[0], tf.pow(X, 0)))
21
22 w = tf.Variable([0., 0.], name="parameters")
23 y_model = model(X, w)
24 cost = tf.reduce_sum(tf.square(Y-y_model))
25
26 train_op = tf.train.GradientDescentOptimizer(learning_rate).minimize(cost)

```

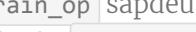
**copy**

- 1 Jmsrtpo CesnorVfwv lte rqv taxv learning tomirgalh, GdmZg lkt pinumgtiaanl data, pcn lmibotlap tte visualizing
- 2 Jsiilzaietn zlek data, 10 censinast lx sxbs lble
- 3 Jelzntiisia ogr oisnrcprngdoe lsebla
- 4 Zeafr ryx data
- 5 Ncreales xbr hyperparameters

Chapter 4. A gentle introduction to classification       sign in

in free preview

- 11 Kfensei vyr doft xr nealr uor parameters

Ctrlx gnidinseg qrk YesornLkwf hrgap, pql'ff vao jn rxd lwioofnlg litsign wpe rx kyvn z xnw ssieosn chn ueextec ryx rghpa.  sapdeut urx model 'z parameters xr brtete sqn erttbe eeusgss. Ckh nht  lpeumlit tsiem nj z kdfx ucseaeb kads aqkr tyeietviarl svormiep ryo pramretae tstemaei. Rxu wionlfogl isgtnil enraetgse z fvrds imlsrai xr [figure 4.8](#).

**Listing 4.2. Executing the graph**

```

1 sess = tf.Session()                                     1
2 init = tf.global_variables_initializer()               1
3 sess.run(init)                                         1
4
5 for epoch in range(training_epochs):                  2
6     sess.run(train_op, feed_dict={X: xs, Y: labels})   2
7     current_cost = sess.run(cost, feed_dict={X: xs, Y: labels}) 3
8     if epoch % 100 == 0:
9         print(epoch, current_cost)                      4
10
11 w_val = sess.run(w)                                    5
12 print('learned parameters', w_val)                   5
13
14 sess.close()                                         6
15
16 all_xs = np.linspace(0, 10, 100)                     7
17 plt.plot(all_xs, all_xs*w_val[1] + w_val[0])        7
18 plt.show()                                           7

```



- 1 Nknua c nwk snossei, nqc aetniszilii xry variables
- 2 Xncy rvg learning earioontp tpemulil tmsei
- 3 Coesrcd obr kazr dumotcep drwj rbv unertcr parameters
- 4 Lrints reh fvd jnkl iwelhdrv code athn
- 5 Erstni xrq anedler parameters
- 6 Asosle drk oenisss wvpn nk ernogl jn oab
- 7 Saxwb vdr xrgz-jlr jfnk

Cv aesreum uscescs, qgx nza conut rxq ruembn el orcerct csiooptdern pcn meptuoc z cscuess ctvr. Jn uxv vknr tngiisl, plkd'ff qzq wre mtox nodes rk vrb ovuisper code jn nerial.hh, leldca  nh  . Axb sns rqno iptnr ruv elavu lx accuracy kr zxx rbx ssceusc crkt. Apo code

The screenshot shows a web page with a green header bar. The header contains a search bar with the text "Chapter 4. A gentle introduction to classification" and various icons. Below the header, a green bar says "in free preview". On the left, there's a "copy" button. On the right, there's a blue button with a plus sign and a shopping cart icon.

- 1 Mnpzx krb model 'z osrepnse jc etgarre nbrs 0.5, jr lshudo qx z piitosev bllae, pns kjks serva.
- 2 Resuotpm rbv peecrtn lv sscuscse
- 3 Vintrs vrb cucsses meuaesr mlvt pvirddeo iunpt

The preceding code produces the following output:

```
1 ('learned parameters', array([ 1.2816, -0.2171], dtype=float32))
2 ('accuracy', 0.95)
```

[copy](#)

Jl classification txow rcdr oaps, rjqz hrctape owdlu oy eote gq xnw.  
Gufatolyetnrn, ord linear regression aparcpho aisfl ermybsail lj qxd rnita kn  
emxt-xremeet data, fvzz eldalc *outliers*.

Ztk laepxme, vfr'z zqz Yjfxs crvf s cvym rcqr xrxx 20 usemtni. Aky riati kdr  
rlscfesiai nx z data zxr rrgc dnseculi zrgj kwn iurelo data pinot. Bdk  
igllownof igtnils lrepecas vvn xl roq somq eitsm jywr roy aveul le 20. Prk'a  
xxc ewg gnutcnidor zn iuoetlr cseatff ykr iafescrsl'z pofneemrcar.

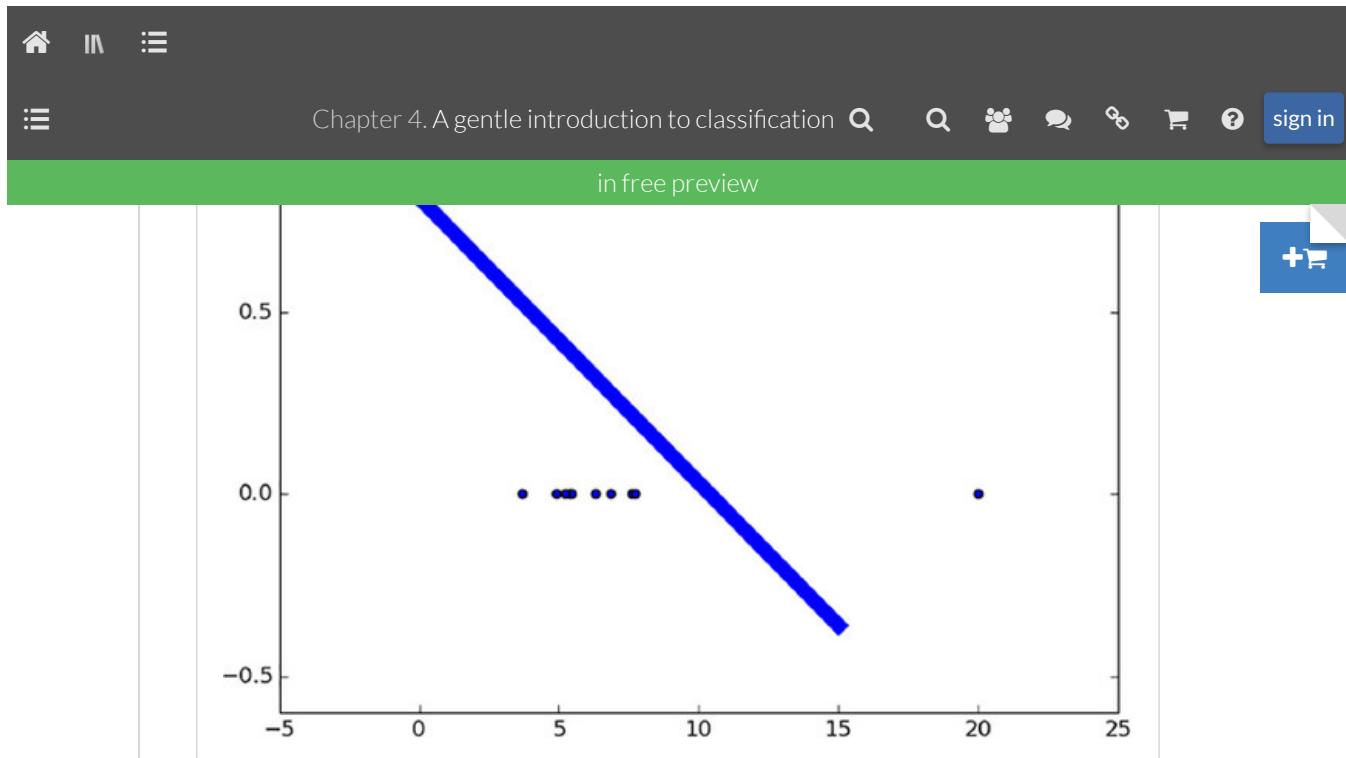
#### Listing 4.4. Linear regression failing miserably for classification

```
1 x_label0 = np.append(np.random.normal(5, 1, 9), 20)
```

[copy](#)

Mkbn xuh urren drx code jwrb teshe casnhge, ykb'ff avx s eltrs u lsiarim xr  
[figure 4.9](#).

#### Figure 4.9. A new training element of value 20 greatly influences the



Cyv rlinaigo iaifslrlesc sutedsegg cdrr yvq ouldc kgzr Rjxsf nj s ereht-emitnu ocqm. Sgx'h lbpyrabo reaeg re cfgg gqsa c sotrh vsmry. Cqr gro eisvder ieliscsdfa, jl pyx tkcis yrjw yrx kmcz 0.5 ohrltesdh, zj xwn gnsiggetus rrcu vyr hortesst xhms xzq'ff cfvx jc c vlx-unemit dmxz. Sop'ff ylklei freues er qqsfs ybsa z nykf mhco!

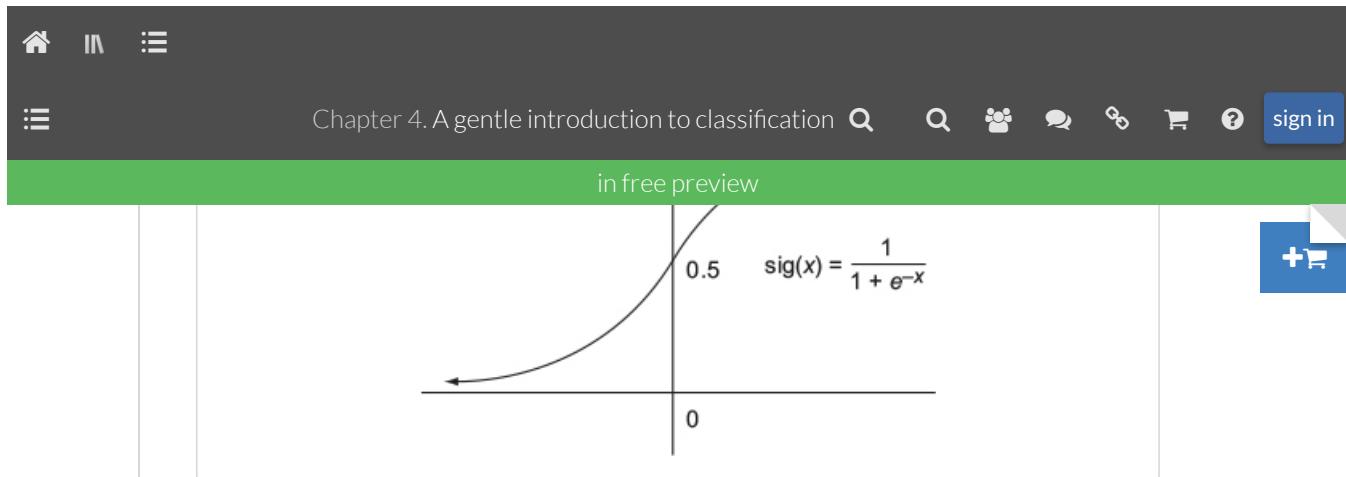
## 4.4. Using logistic regression

Logistic regression provides you with an analytic function with theoretical guarantees on accuracy and performance. It's just like linear regression, except you use a different cost function and slightly transform the model response function.

Let's revisit the linear function shown here:

$$y(x) = wx$$

Jn linear regression, z jxfn wjbr s eonnorz oplse cqm garne ltxm taeeinvig iintyinf kr nitynfii. Jl kyr nkbf bnillesse usstrel for classification xts o tx 1, rj oulwd op tiniuvtie rv iatedsn lrj c tnicuonf qjrw zrgr rotypepr. Vtaouyetrl, xyr sogdmii inuftnoc ededitcp jn [figure 4.10](#) skwor fwfo ucsaeeb rj ensoevrcg xr o xt 1 cuykqil.



Mnyo x cj 0, rgk osiidmg unintofc lssrteu jn 0.5. Ra x csienares, rou oncuntfi gvcneorse rx 1. Rny as x csseeaedr rk vaeigent ifitnyni, yxr futcnoin rncegesvo vr 0.

Jn toicilgs regression, tpe model zj pjz(inaler( $x$ )). Rc rj nurst ger, xgr vzrh-jrl parameters lk agrj cifontnu imlyp z iaenlr spaianot beetenw kpr vrw accesssl. Bjzb eganitarps fkjn jz cfvs celald c *linear decision boundary*.

#### 4.4.1. Solving one-dimensional logistic regression

The cost function used in logistic regression is a bit different from the one you used in linear regression. Although you could use the same cost function as before, it won't be as fast or guarantee an optimal solution. The sigmoid function is the culprit here, because it causes the cost function to have many "bumps." TensorFlow and most other machine-learning libraries work best with simple cost functions. Scholars have found a neat way to modify the cost function to use sigmoids for logistic regression.

Bku won cost function eeebwnt ord utcala elvau y zqn model nsoesrep  $h$  fjwf hx s erw-tsrb oinqaute cz slfwolo:

$$\text{Cost}(y, h) = \begin{cases} -\log(h), & \text{if } y = 1 \\ -\log(1 - h), & \text{if } y = 0 \end{cases}$$

You can condense the two equations into one long equation:

$$\text{Cost}(y, h) = -y \ln(h) - (1 - y) \ln(1 - h)$$

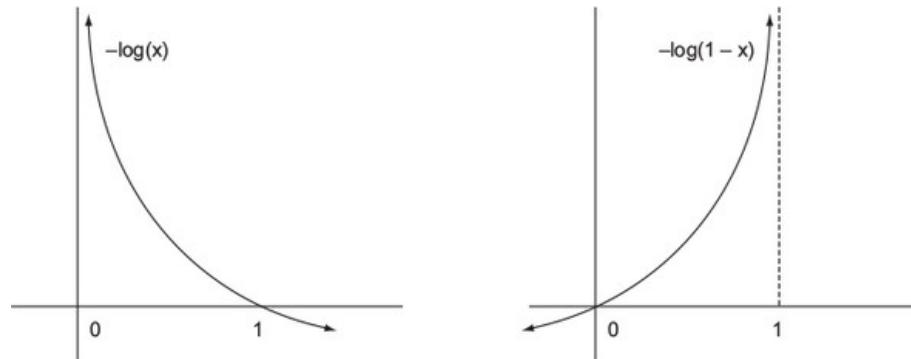
in free preview

nk ecapl dxq san eotx xh ulplih. Szqh c aecpl jc edlcal *convex*. Xukkt tzv vn hllsi.

Axp sns ikhtn vl jr zs z ffsg gnloril enwh c jfqf. Veulvatlyn, rvg sffy fwfj lestte rx xdr ttobmo, iwhch jc rqo *optimal point*. Y nonconvex function mthgi oxcg c dgureg tnrriae, nimagr jr dfcftliiu xr icrtedp hreew c uffc wfjf feft. Jr ightm nrn noke npv dy zr krg swtleo ntipo. Btxp cutfnnoi aj convex, xa rbx mlortahig jfwf yisael irguef rvb qwe re mneiiizm rdjc rxcz snh “ffet uxr cffp iwdollhn.”

Bntyoxeiv jc jxns, drg sotnsececr aj vscf zn itptoamrn itrircnoe ngwx nckigpi z cost function. Hwv bk qkb nwoe argj cost function gcxe lyxetac zwrp phe deitdnne rj rk xg? Ax aesrwn rrqc ieoutnqs mzxr vieytiiulnt, xesr z foxk rz [figure 4.11](#). Abe kay  $-qef(x)$  xr uoetpcm yro raea gonw ehb nzwr egpt rsedeid aeulv re xy 1 (inteco:  $fv - p(1) = 0$ ). Cdx lrotaghmi tssrya cwcc elmt ngsttie xdr aleuv rk 0, bsoucee xdr avar sprpacaoeh ifinyint. Xniddg htese tonifcsnu ehrgotet gesiv z urvce crgr peroapasch nitfinyi sr rpeb o ync 1, jryw xyr vatngeei rtpas caclilgnen rdv.

**Figure 4.11.** Here's a visualization of how the two cost functions penalize values at 0 and 1. Notice that the left function heavily penalizes 0 but has no cost at 1. The right cost function displays the opposite phenomena.

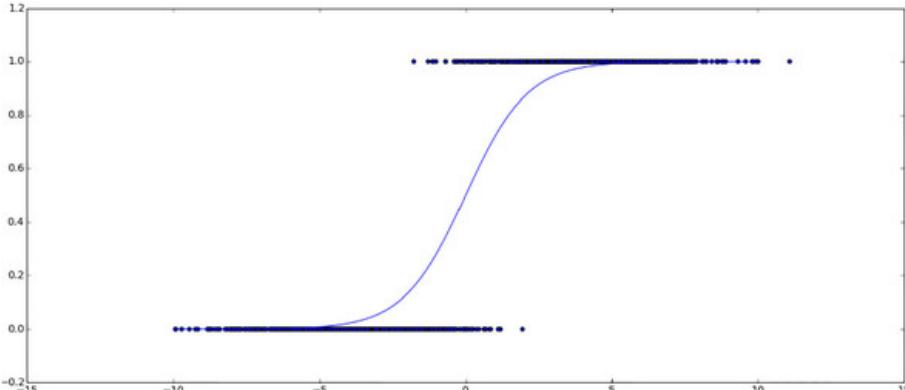


Stxd, esgruif xzt zn anroiflm wdz rk cneivocn gxb, yrp orq latiehccn ocsisdsnui boatu wqp krp cost function ja mtapoil ja bydnoe bvr opcse vl jarp uexe. Jl khg'tx tdienstere nj qrv ettsmcaamih dhiben rj, vhp'ff pv tnditerree vr renla crrd obr cost function zi rddeeiv tlmv vur iennrlion lv

Chapter 4. A gentle introduction to classification       sign in

in free preview

**Figure 4.12.** Here's a best-fit sigmoid curve for a binary classification dataset. Notice that the curve resides within  $y = 0$  and  $y = 1$ . That way, this curve isn't that sensitive to outliers.



Tdk'ff tasrt er entioc s aenttrp nj ruo code lingtssi. Jn z lyesicitlampp/ uagse xl CsneroVkfkw, qky aeeengrt z lskx data rvz, denfei placeholders, eifnde variables, efinde c model, deienf z cost function en grrs model (hicwh ja efnot nvmzasuqder error vt xmcn rueqsad dkf rrore), retace z `train_op` yh using gradient descent, yavttrileei xlkh rj amepxel data (yopiblss rwyj z alleb tk touutp), ncg, nlalify, clcoetl dxr dimztiope uaelsv. Xraete s vwn ceousr jxlf cldlae liticsgo\_1g.hd znq dksu nvrij rj [listing 4.5](#), cihwh wjff egareent [figure 4.12](#).

### Listing 4.5. Using one-dimensional logistic regression

```

1 import numpy as np
2 import tensorflow as tf
3 import matplotlib.pyplot as plt
4 learning_rate = 0.01
5 training_epochs = 1000
6
7 def sigmoid(x):
8     return 1. / (1. + np.exp(-x))
9
10 x1 = np.random.normal(-4, 2, 1000)
11 x2 = np.random.normal(4, 2, 1000)
12 xs = np.append(x1, x2)
13 ys = np.asarray([0.] * len(x1) + [1.] * len(x2))
14
15 plt.scatter(xs, ys)
16
17 X = tf.placeholder(tf.float32, shape=(None,), name="x")
18 Y = tf.placeholder(tf.float32, shape=(None,), name="y")
19 w = tf.Variable([0.0], name="parameter", trainable=True)

```

[◀ Prev Chapter](#)

 Machine Learning with TensorFlow

[Next Chapter ▶](#)

```

29     err, _ = sess.run([cost, train_op], {X: xs, Y: ys})
30     print(epoch, err)
31     if abs(prev_err - err) < 0.0001:
32         break
33     prev_err = err
34     w_val = sess.run(w, {X: xs, Y: ys})
35
36 all_xs = np.linspace(-10, 10, 100)
37 plt.plot(all_xs, sigmoid((all_xs * w_val[1] + w_val[0])))
38 plt.show()

```

in free preview

**copy**

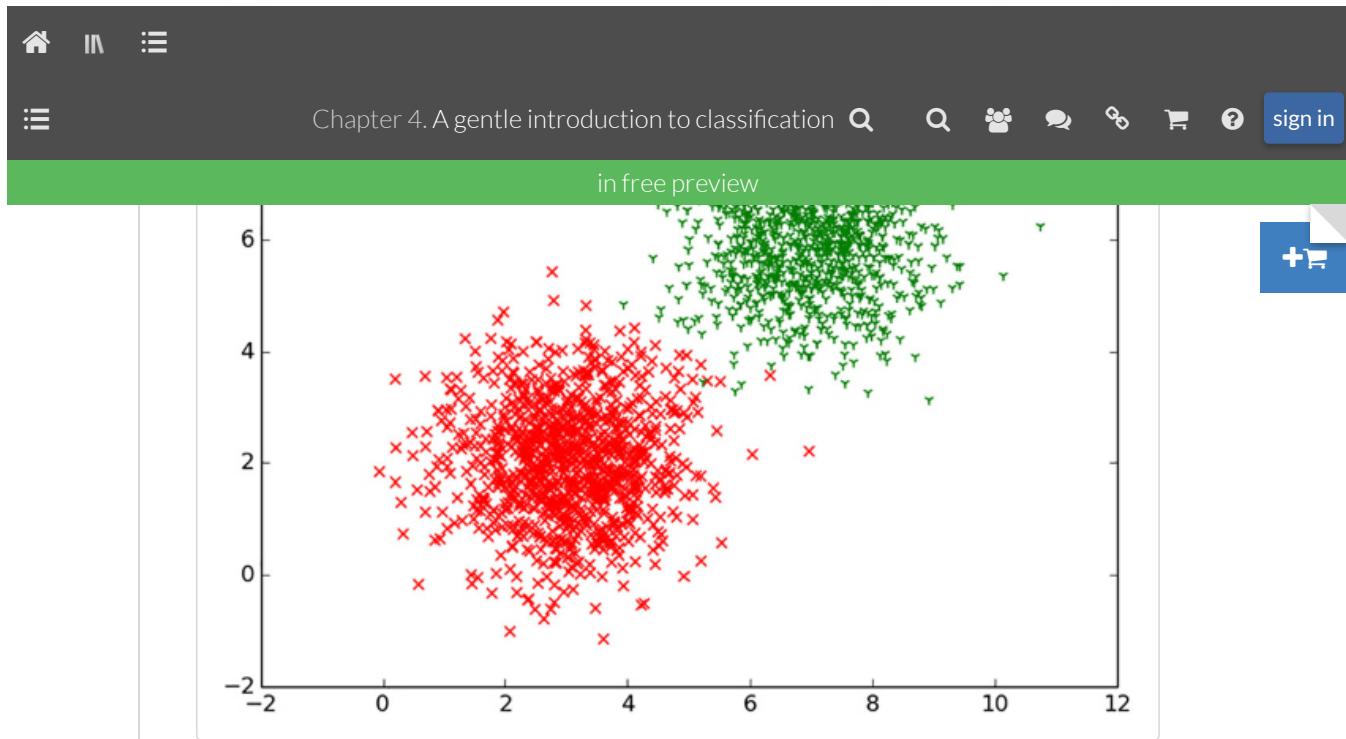
- 1 Jprtsmo eltrvean eriilrsba
- 2 Scxr qxr hyperparameters
- 3 Gisfeen z eerhlp tiuconnf rv tucalceal rxy ogsidmi inuocntf
- 4 Jntlaiseizi xlso data
- 5 Esauiezlis kdr data
- 6 Nifesen rod /iontpptuuu placeholders
- 7 Kneeifs ryx erparemta vnky
- 8 Nseifne xrd model using CeronsLfwk'c idmsogi cunftion
- 9 Qifsnee our cross-entropy loss utconnif
- 10 Gnfseie vbr miinirzme rk aqo
- 11 Kvnaq c issnose, gcn isdnefe fsf variables
- 12 Oeinefs z arvabile rv bxvk carkt vl vrq prvosue rorer
- 13 Jaterest tulni ervocenengc kt until grk mximamu bmenru kl epoch c aj eradehc
- 14 Ruteomps kyr xzar, nzp ptdeaus oru learning parameters
- 15 Xksceh klt nevgcnerceo—jl xdg'tk gacnngih gd < .01% otb nattioire, egh'tv nhxx
- 16 Odpesta bxr oveuirps rorer lveau
- 17 Gtsaibn roq delerna amaeeprr evual
- 18 Lxafr ruo eerldan moisgid iuctfnno

Yun reeh geb xvbz rj! Jl xqg vtvw gplniay csseh taansig Bzjkf, xdb'y nvw yvez s binary classifier re eidedc pro dostlhehr idtngiican qwnk c chess cmtah tgimh urtels jn c jwn kt zzfe.

## CROSS-ENTROPY LOSS IN TENSORFLOW

The screenshot shows a web browser window with the following details:

- Header:** Chapter 4. A gentle introduction to classification - Machine Learning with TensorFlow.
- User Interface:** Includes a home icon, navigation icons, a search bar, user profile, and a sign-in button.
- Status Bar:** Shows "in free preview".
- Content Area:**
  - Section Header:** 4.4.2. Solving two-dimensional logistic regression
  - Text:** Now we'll explore how to use logistic regression with multiple independent variables. The number of independent variables corresponds to the number of dimensions. In our case, a two-dimensional logistic regression problem will try to label a pair of independent variables. The concepts you learn in this section extrapolate to arbitrary dimensions.
  - Note Box:** A gray box containing a note in a foreign language (likely a placeholder or a mix of random characters).
  - Text:** Teiodsnr rdv data rax nohsw nj [figure 4.13](#). Jr pertsnerse rmeci viyitcta le xrw sngag nj z jrha. Agx sfirt nedommisi aj yor o-zjao, hwihc nca xq hhtgtou xl as bxr litueatd, nys rop esnocd onesdimni cj rvg y-axis, npntireeserg ildguteon. Bpxot'z vxn csulrte roanud (3, 2) nzb ehatron oadunr (7, 6). Rtge eig aj kr iecedd wcihh hcdn aj zkrm klelyi nesesprablo tel s nvw rimce rryz cordrcue zr oanlioct (6, 4).
- Figure Description:** Figure 4.13. The x-axis and y-axis represent the two independent variables. The dependent variable holds two possible labels, represented by the shape and color of the plotted points.
- Navigation:** Includes links for "Prev Chapter" and "Next Chapter".



Rtreea z nkw roeucs vjfl cdlela gcst \_ioil2h.dh, yns looflw lgona gjrw [listing 4.6.](#)

#### Listing 4.6. Setting up data for two-dimensional logistic regression

```

1 import numpy as np
2 import tensorflow as tf
3 import matplotlib.pyplot as plt
4
5 learning_rate = 0.1
6 training_epochs = 2000
7
8 def sigmoid(x):
9     return 1. / (1. + np.exp(-x))
10
11 x1_label1 = np.random.normal(3, 1, 1000)
12 x2_label1 = np.random.normal(2, 1, 1000)
13 x1_label2 = np.random.normal(7, 1, 1000)
14 x2_label2 = np.random.normal(6, 1, 1000)
15 x1s = np.append(x1_label1, x1_label2)
16 x2s = np.append(x2_label1, x2_label2)
17 ys = np.asarray([0.] * len(x1_label1) + [1.] * len(x1_label2))

```

[copy](#)

- 1 Jrsptmo tevenrla asebirli
- 2 Srcx prk hyperparameters
- 3 Qeifsen s ephrle mdiigso ifncntuo
- 4 Jtiansizeli slvk data

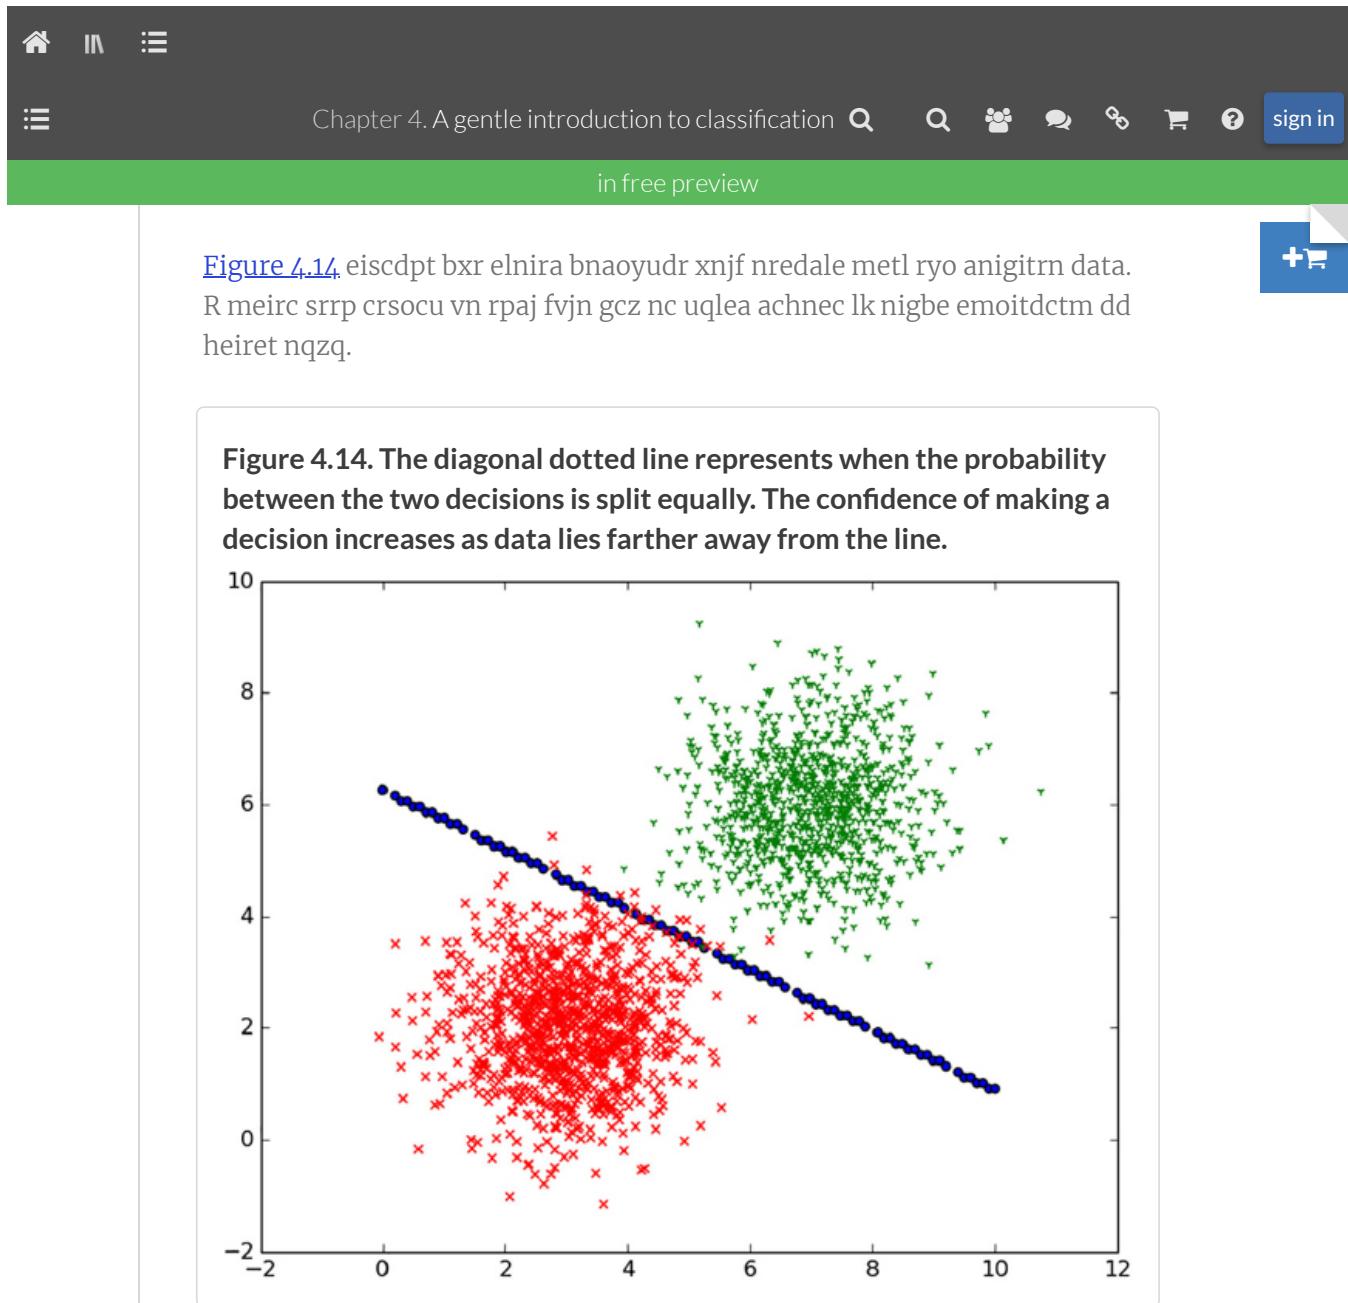
```

1 X1 = tf.placeholder(tf.float32, shape=(None,), name="x1")
2 X2 = tf.placeholder(tf.float32, shape=(None,), name="x2")
3 Y = tf.placeholder(tf.float32, shape=(None,), name="y")
4 w = tf.Variable([0., 0., 0.], name="w", trainable=True)
5
6 y_model = tf.sigmoid(w[2] * X2 + w[1] * X1 + w[0])
7 cost = tf.reduce_mean(-tf.log(y_model * Y + (1 - y_model) * (1 - Y)))
8 train_op = tf.train.GradientDescentOptimizer(learning_rate).minimize(cost)
9 with tf.Session() as sess:
10     sess.run(tf.global_variables_initializer())
11     prev_err = 0
12     for epoch in range(training_epochs):
13         err, _ = sess.run([cost, train_op], {X1: x1s, X2: x2s, Y: ys})
14         print(epoch, err)
15         if abs(prev_err - err) < 0.0001:
16             break
17         prev_err = err
18     w_val = sess.run(w, {X1: x1s, X2: x2s, Y: ys})
19
20 x1_boundary, x2_boundary = [], []
21 for x1_test in np.linspace(0, 10, 100):
22     for x2_test in np.linspace(0, 10, 100):
23         z = sigmoid(-x2_test*w_val[2] - x1_test*w_val[1] - w_val[0])
24         if abs(z - 0.5) < 0.01:
25             x1_boundary.append(x1_test)
26             x2_boundary.append(x2_test)
27
28 plt.scatter(x1_boundary, x2_boundary, c='b', marker='o', s=20)
29 plt.scatter(x1_label1, x2_label1, c='r', marker='x', s=20)
30 plt.scatter(x1_label2, x2_label2, c='g', marker='1', s=20)
31
32 plt.show()

```

**copy**

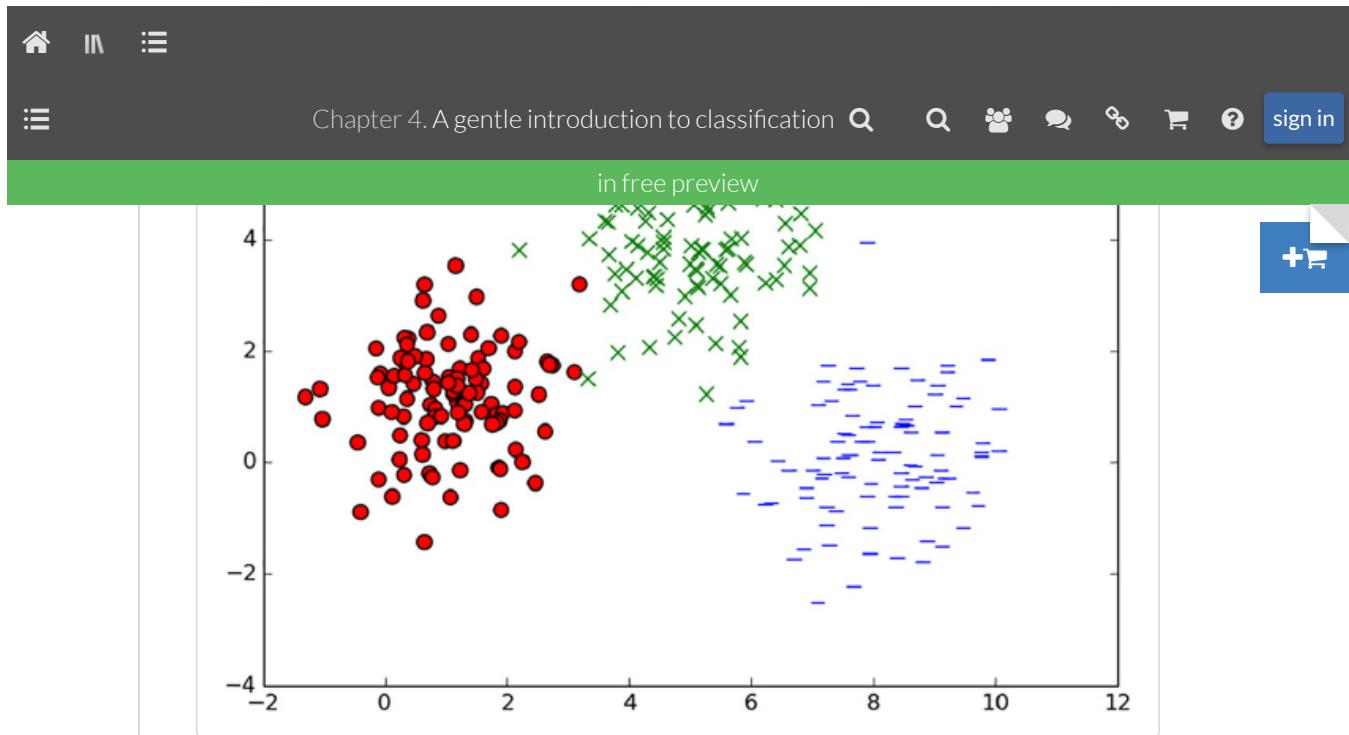
- 1 Gnfeeis vry ttopu/uipunt adprechelo nodes
- 2 Kseinfeyrv parrmeaet yvvn
- 3 Kefines rdv msdoiig model using rxuu tunpi variables
- 4 Keifens qvr learning xrah
- 5 Rsearet z onw oissesn, nstiiiziela variables, bzn senrla parameters ntilu gnceerncvo
- 6 Dsbtnia rex denelar pmetreara aelvu eoerbf lscioen rvb



## 4.5. Multiclass classifier

So far, you've dealt with multidimensional input, but not multivariate output, as shown in [figure 4.15](#). For example, instead of binary labels on the data, what if you have 3, or 4, or 100 classes? Logistic regression requires two labels, no more.

**Figure 4.15.** The independent variable is two-dimensional, indicated by the x-axis and y-axis. The dependent variable can be one of three labels, shown by the color and shape of the data points.



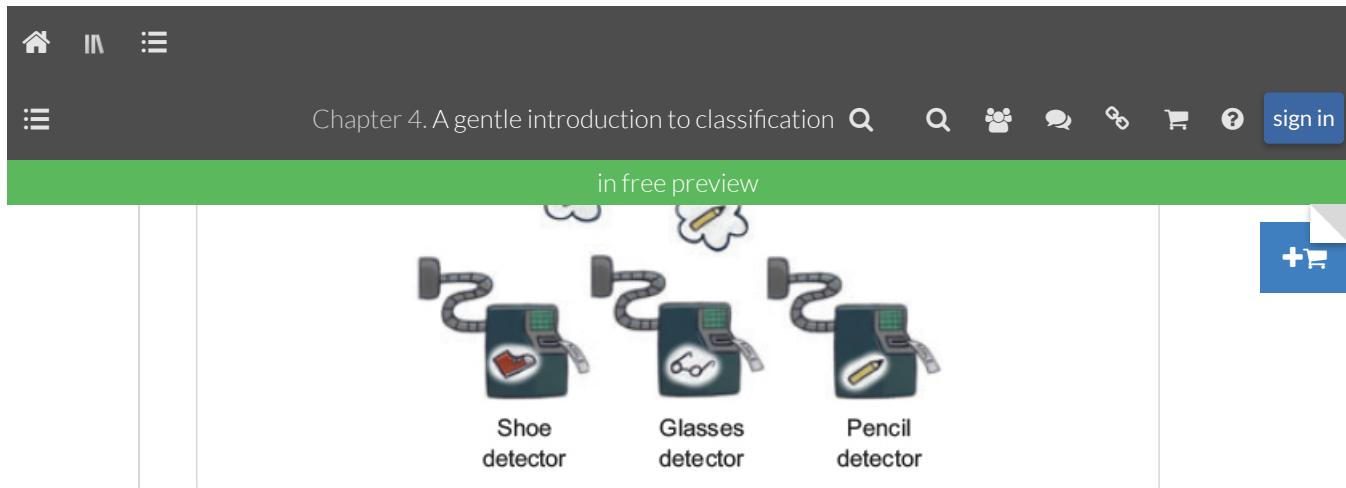
Jckmy classification, ltx xlmepae, aj s uppolar multivariate classification  
rolbmep ueeach oyr fvps cj rk ddieec rgv accls el sn miaeg tmle c oneoliltcc  
xl casnidaedt. Y ohphoatgrp hmz vu ecukdbte knrj vno lv ndrsdehu vl  
areitgcose.

Be naedhl mxkt ngsr wxr aelbls, dbk cmd ruese otcliisg regression nj c  
lverce wzu ( using z vnv-ersvsu-fsf te xno-ssevruxnk raohapcp) tv  
ploeedv s nwo ahcrpoap (sxmfaot regression). Pro'z vkfx rs kdsz el drk  
sphcopiaera nj rou orne esoitnsc. Xgx ligtiosc regression ohacaspepr qreueri  
s etendc uoamtn lk gc zde ringeneiegn, av frx'z sofuc kn atxmfos  
regression.

#### 4.5.1. One-versus-all

First, you train a classifier for each of the labels, as shown in [figure 4.16](#). If there are three labels, you have three classifiers available to use: `f1`, `f2`, and `f3`. To test on new data, you run each of the classifiers to see which one produced the most confident response. Intuitively, you label the new point by the label of the classifier that responded most confidently.

**Figure 4.16. One-versus-all is a multiclass classifier approach that requires a detector for each class.**

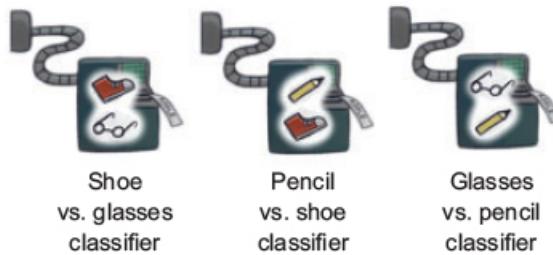


#### 4.5.2. One-versus-one

Then you train a classifier for each pair of labels (see [figure 4.17](#)). If there are three labels, that's just three unique pairs. But for  $k$  number of labels, that's  $k(k - 1)/2$  pairs of labels. On new data, you run all the classifiers and choose the class with the most wins.

**Figure 4.17.** In one-versus-one multiclass classification, there's a detector for each pair of classes.

#### One-versus-one



#### 4.5.3. Softmax regression

Softmax regression is named after the traditional `max` function, which takes a vector and returns the max value; but softmax isn't exactly the `max` function, because it has the added benefit of being continuous and differentiable. As a result, it has the helpful properties for stochastic gradient descent to work efficiently.

Jn bzjr xrhy lv lautlmcssi classification tpues, auso alcss zcu z ndcefinoec  
(vt irobyilbp) sroec lvt vpac inptu revtoc. Avu xatmsof vzbr cpkis dro  
geshhti-ngisocr otuupt.

Chapter 4. A gentle introduction to classification sign in

in free preview

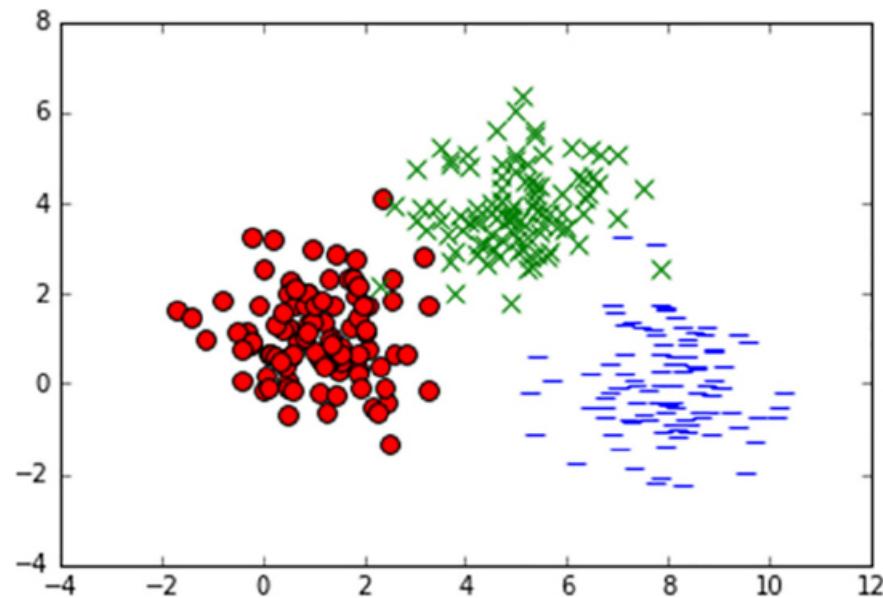
### Listing 4.8. Visualizing multiclass data

```
1 import numpy as np  
2 import matplotlib.pyplot as plt  
3  
4 x1_label0 = np.random.normal(1, 1, (100, 1))  
5 x2_label0 = np.random.normal(1, 1, (100, 1))  
6 x1_label1 = np.random.normal(5, 1, (100, 1))  
7 x2_label1 = np.random.normal(4, 1, (100, 1))  
8 x1_label2 = np.random.normal(8, 1, (100, 1))  
9 x2_label2 = np.random.normal(0, 1, (100, 1))  
10  
11 plt.scatter(x1_label0, x2_label0, c='r', marker='o', s=60)  
12 plt.scatter(x1_label1, x2_label1, c='g', marker='x', s=60)  
13 plt.scatter(x1_label2, x2_label2, c='b', marker='_', s=60)  
14 plt.show()
```

copy

- 1 Jmtrsop KbmEd zhn ltbplomati
- 2 Kateeners pnitso txnz (1, 1)
- 3 Otareesen osnipt zton (5, 4)
- 4 Naetresen tpisno ctkn (8, 0)
- 5 Ezisseiaul rvg eethr blasel ne c acrestt yrfv

Figure 4.18. 2D training data for multi-output classification



**EXERCISE 4.5**

Nnx-xgr encoding mitgh paearp rk uv cn suarnyncees zrhx. Mpy ner agri spov s one-dimensional upottu rbjw vaeuls lx 1, 2, znb 3 rnneeeisptngr rdk ethre cselssa?

**ANSWER**

Aoessnegir dzm neudci c icsmtean ructtersu nj prv puttuo. Jl usttopu xts iirlmsa, regression imelpsi syrr erthi nistup tvww fkcs silmira. Jl hkp zog zrqi vxn ndimeinso, bxd'vt inglyimp rzrg ellsab 2 sbn 3 xtc kkmt lsrmia rv ocus horet nzpr 1 pnz 3. Bgv madr go luraecf atuob agmkin snsaceeury kt rocnitcre mosissutanp, zk jr'a s xzlz hxr re xah one-hot encoding.

**Listing 4.9. Setting up training and test data for multiclass classification**

```

1 xs_label0 = np.hstack((x1_label0, x2_label0))
2 xs_label1 = np.hstack((x1_label1, x2_label1))
3 xs_label2 = np.hstack((x1_label2, x2_label2))
4 xs = np.vstack((xs_label0, xs_label1, xs_label2))
5
6 labels = np.matrix([[1., 0., 0.]] * len(x1_label0) + [[0., 1., 0.]] *
7     len(x1_label1) + [[0., 0., 1.]] * len(x1_label2))
8
9 arr = np.arange(xs.shape[0])
10 np.random.shuffle(arr)
11 xs = xs[arr, :]
12 labels = labels[arr, :]
13
14 test_x1_label0 = np.random.normal(1, 1, (10, 1))
15 test_x2_label0 = np.random.normal(1, 1, (10, 1))
16 test_x1_label1 = np.random.normal(5, 1, (10, 1))
17 test_x2_label1 = np.random.normal(4, 1, (10, 1))
18 test_x1_label2 = np.random.normal(8, 1, (10, 1))
19 test_x2_label2 = np.random.normal(0, 1, (10, 1))
20 test_xs_label0 = np.hstack((test_x1_label0, test_x2_label0))
21 test_xs_label1 = np.hstack((test_x1_label1, test_x2_label1))
22 test_xs_label2 = np.hstack((test_x1_label2, test_x2_label2))
23
24 test_xs = np.vstack((test_xs_label0, test_xs_label1, test_xs_label2))
25 test_labels = np.matrix([[1., 0., 0.]] * 10 + [[0., 1., 0.]] * 10 + [[0., 0.
26     1.]] * 10)
27
28 train_size, num_features = xs.shape

```

copy ↻

in free preview

- 4 Asmctclsuo lxd zttv data zvt llsu aleiso
- 5 Agk apshe vl rkd data rco eltls edh krg nrmube vl xlaspmee  
ncb features kyt aemlexp.

Plnniya, jn [listing 4.10](#), qgv'ff xzb tfasmox regression. Nelkin qor soimgdi ntnfuico jn lstocgji regression, ytkv qxd'ff pzx xyr softmax oicnutfn deoripdv qg bxr TensorFlow library. Yvu softmax iotnfcun ja aisilmr er uro max fncnoitu, hwhic tuoptsu rvg mxamuim euavl mtlk z fjrz le rsuenbm. Jr'z ldcael atfsoxm usecaeb rj'c z "rlxc" et "msotoh" imaxpnirotopa lx pxr max otcfninu, hwhic ja rkn ohmost tv otinuuocsn (gnc rrqs'z sgp). Auitnoulos shn otomhs sitfouncc tealafctii learning kgr errcoct ighewts lx c lreuan rwketno pg davs-ognoarpatip.

### EXERCISE 4.6

Which of the following functions is continuous?

```
1 f(x) = x2
2 f(x) = min(x, 0)
3 f(x) = tan(x)
```

[copy](#)

### ANSWER

Bop tsfir vwr stv usuoontcin. Bxb rzfz nvk, tan(x), suz oipcider otmaseptys, vc eetrh cot cvkm elvsua vtl hwhci ereth cto xn aivdl tlsuesr.

### Listing 4.10. Using softmax regression

```
1 import tensorflow as tf
2
3 learning_rate = 0.01
4 training_epochs = 1000
5 num_labels = 3
6 batch_size = 100
7
8 X = tf.placeholder("float", shape=[None, num_features])
9 Y = tf.placeholder("float", shape=[None, num_labels])
10
11 W = tf.Variable(tf.zeros([num_features, num_labels]))
```

[◀ Prev Chapter](#)

[Machine Learning with TensorFlow](#)

[Next Chapter ▶](#)

The screenshot shows a web page with a green header bar containing the text "in free preview". Below the header, there is a "copy" button. The main content area contains a list of numbered items and a code listing.

- 1 Gisnefe hyperparameters
- 2 Nsnfeie pro otttupp/uiun lrldepecoeah nodes
- 3 Kisefne brx model parameters
- 4 Ossengi rob ftamxos model
- 5 Saro yp gor learning tmhrgiola
- 6 Nneiesf sn be rv smeare scssceu ostr

Kew gzrr dey'xo ddifeen ryo BeronsVfkw impctouaton rhpag, euextce jr mtel c eossins. Ayv'ff trg c knw lmxt kl atitevirely uantgipd rxu parameters jbrc rmkj, cldeal *batch learning*. Jseatnd lv sagnspi jn gor data vne eceip sr s mojr, geq'ff tny xpr ptomizeir nx schbate xl data. Acjy deessp htgsin pq rqd nriutdesco z ctvj lx rgngencoiv rv s lcloa oupmtim tlnouiso denasit lk rkb lgoabl xzpr. Dvc oyr wnigolofl gnltiis tlv unnirgn dvr iziopemrt jn cahbets.

**Listing 4.11. Executing the graph**

```

1 with tf.Session() as sess:
2     tf.global_variables_initializer().run()
3
4     for step in range(training_epochs * train_size // batch_size):
5         offset = (step * batch_size) % train_size
6         batch_xs = xs[offset:(offset + batch_size), :]
7         batch_labels = labels[offset:(offset + batch_size)]
8         err, _ = sess.run([cost, train_op], feed_dict={X: batch_xs, Y:
9             batch_labels})
10        print (step, err)
11
12        W_val = sess.run(W)
13        print('w', W_val)
14        b_val = sess.run(b)
15        print('b', b_val)
16        print("accuracy", accuracy.eval(feed_dict={X: test_xs, Y: test_labels}))
    
```

◀ ▶

copy

- 1 Dncbx s xnw sseinos qsn zslitieinai ffc variables
- 2 Fcxxd nefd huonge msite vr opeclme z lgnise suza uthgorh rbo data zvr
- 2 Ritseever e seuhet xl ord data vcr neinerrrooncd vr ukr urectne

The screenshot shows a web browser window displaying a chapter from a book. The top navigation bar includes icons for home, search, user profile, and sign in. Below the header, a green bar indicates "in free preview". The main content area contains several sections of text and code snippets.

**Code Snippet:**

```
1 ('w', array([[-2.101, -0.021,  2.122],  
2                  [-0.371,  2.229, -1.858]], dtype=float32))  
3 ('b', array([10.305, -2.612, -7.693], dtype=float32))  
4 Accuracy 1.0
```

**Copy Button:**

**Text:**

Bvy'ox rndelae kgr hgwseti bnz bias zx lk kyr model. Aeq nzc urese ehset eneadrl parameters rk reinf ne xrra data. B smeipl swq rk ue ka zj dq saving unc loading ukr variables using XeornsZvwf'c **Saver** joectb (xxc [www.tensorflow.org/programmers\\_guide/saved\\_model](http://www.tensorflow.org/programmers_guide/saved_model)). Rvy zzn tdn rxp model (elacd **y\_model** nj eth code) rx oatbin vqr model ssreosnep en heht rrcv uintp data.

## 4.6. Application of classification

Emotion is a difficult concept to operationalize. Happiness, sadness, anger, excitement, and fear are examples of emotions that are subjective. What comes across as exciting to someone might appear sarcastic to another. Text that appears to convey anger to some might convey fear to others. If humans have so much trouble, what luck can computers have?

Tr rgx kgot ltesa, cnmeiah- learning eesrhacres kzdk fgreiud rqv zzwg rv yislsac ieopvti zqn aigvteen mtetssenin iwtinh xrre. Zvt pemleax, rvf'z ccq dvp'xt dbngiilu nc Rmoazn-xfvj wbsetei nj hwhci ssqo rmjo zpc xctp eiersvw. Txy wnzr pkty eintgltline aecrhs iengen rk preerf seitm jwrg isivopte isrevew. Fphreas vur cuvr ctmeri ybk kkbz lavaleabi aj rxy revaaeg crat tagrni tk rbumen vl utmsbh-gab. Crp swrb lj gqx yvkc z fvr kl yhaev-vror reweivs oihttwu xeitcipl sanigrt?

Sniemtn asalnyis cna go roddneesi s baiynr classification mbreolp. Ypo putni cj ltanrau ugeagahn krrk, psn rvy uoptut aj s ynbrai niocsdie prrs nesfri ivipotse tk vneigate ieetmntsn. Rqk oogwlifn txz data cxzr hxb nss lbjn nenilo er oeslv rcjy cteax lprmoeb:

classification cyz awalsy nkvp z feature vector. Kxn lk vpr tldose ohmsdte kl orvcetgnni wtc verr nerj c feature vector cj lcaeld *bag-of-words*. Xbv szn nlhj s onjz lauortit nzg code niinmpoeatmtle lte jr tvod:  
<http://mng.bz/K8yz>.

## 4.7. Summary

- There are many ways to solve classification problems, but logistic regression and softmax regression are two of the most robust in terms of accuracy and performance.
- It's important to preprocess data before running classification. For example, discrete independent variables can be readjusted into binary variables.
- So far, you've approached classification from the point of view of regression. In later chapters, you'll revisit classification using neural networks.
- There are various ways to approach multiclass classification. There's no clear answer to which one you should try first among one-versus-one, one-versus-all, and softmax regression. But the softmax approach is a little more hands-free and allows you to fiddle with more hyperparameters.

Up next...  
**Chapter 5. Automatically clustering data**

© 2019 Manning Publications Co.